


What type of initialization can all types use?

- Structured binding
 - `auto [first, second] = p;`
- Member-wise
 - `student.name = "Jacob"`
-  **Uniform initialization**
 - `Student jacob { "Jacob", "NM", 21 }`

A **stringstream** is an...

- Input stream
- Output stream
-  **Both!**
- Neither!



The many containers of C++

The C++ Standard Template Library (STL)

`std::vector`

`std::unordered_map`

`std::set`

`std::unordered_set`

`std::stack`

`std::priority_queue`

`std::queue`

`std::deque`

`std::map`

`std::array`


What are templates?

```
class IntVector {  
    class DoubleVector {  
        class StringVector {  
            // Code to store  
            // a list of  
            // strings...  
        };  
    };  
};
```

```
template <typename T>  
class vector {  
    // So satisfying.  
};  
  
vector<int> v1;  
vector<double> v2;  
vector<string> v3;
```

All STL containers are templates!

The Standard Template Library (STL)



The Standard Template Library (STL)

Containers

How do we store groups of things?

Iterators

How do we traverse containers?

Functors

How can we represent functions as objects?

Algorithms

How do we transform and modify containers in a generic way?

Sequence containers store a linear sequence of elements


std::vector

#include <vector>

std::vector stores a list of elements

```
std::vector<int> vec { 1, 2, 3, 4 };
vec.push_back(5);
vec.push_back(6);
vec[1] = 20;


for (size_t i = 0; i < vec.size(); i++) {
    std::cout << vec[i] << " ";
}
```



Stanford vs. STL vector

What you want to do?	Stanford Vector<int>	std::vector<int>
Create an empty vector	Vector<int> v;	std::vector<int> v;
Create a vector with n copies of 0	Vector<int> v(n);	std::vector<int> v(n);
Create a vector with n copies of value k	Vector<int> v(n, k);	std::vector<int> v(n, k);
Add k to the end of the vector	v.add(k);	v.push_back(k);
Clear vector	v.clear();	v.clear();
Check if v is empty	if (v.isEmpty())	if (v.empty())
Get the element at index i	int v = v.get(i); int k = v[i];	int k = v.at(i); int k = v[i];
Replace the element at index i	v.get(i) = k; v[i] = k;	v.at(i) = k; v[i] = k;

How is vector implemented?



size = 5, capacity = 8

Tip: Use range-based for when possible

```
for (size_t i = 0; i < vec.size(); i++) {  
    std::cout << vec[i] << " ";  
}
```



```
for (auto elem : vec) {  
    std::cout << elem << " ";  
}
```

- Applies for all iterable containers, not just `std::vector`

Tip: Use **const auto&** when possible

```
std::vector<MassiveType> vec { ... };  
for (auto elem : vec) ...
```



```
for (const auto& elem : v)
```

- Applies for all iterable containers, not just std::vector
- Saves making a potentially expensive copy of each element

operator[] does not perform bounds checking

```
std::vector<int> vec{5, 6}; // {5, 6}
vec[1] = 3;                // {5, 3}
vec[2] = 4;                // undefined behavior
vec.at(2) = 4;             // Runtime error
```

std::vector is not the best for all cases...


- Suppose we need to observe the last 10,000 prices of a stock
 - What might be concerning about the code below?



Trick question!

std::vector has no push_front!

A hypothetical `push_front`...



SLOW!!!


std::deque
#include <deque>

std::deque

- A deque (“deck”) is a double-ended queue
 - Allows efficient insertion/removal at either end

A **deque** has the same interface as **vector**,
except we can **push_front / pop_front**

How is **deque** implemented?




The problem with **vector** is
that we have a single
chunk of memory

So... let's split it up!

How is deque implemented?

Array of arrays



Separate subarrays allocated independently

Associative containers organize elements by unique keys

std::map
#include <map>


`std::map` maps keys to values



- Equivalent of a Python dictionary
- Sometimes called an **associative array**

`std::map` maps keys to values

```
std::map<std::string, int> map {  
    { "Chris", 2 },  
    { "CS106L", 42 },  
    { "Keith", 14 },  
    { "Nick", 51 },  
    { "Sean", 35 },  
};  
  
int sean = map["Sean"]; // 35  
map["Chris"] = 31;
```



Stanford vs. STL map

What you want to do?	Stanford Map<char, int>	std::map<char, int>
Create an empty map	<code>Map<char, int> m;</code>	<code>std::map<char, int> m;</code>
Add key k with value v into the map	<code>m.put(k, v);</code> <code>m[k] = v;</code>	<code>m.insert({k, v});</code> <code>m[k] = v;</code>
Remove key k from the map	<code>m.remove(k);</code>	<code>m.erase(k);</code>
Check if k is in the map (* C++20)	<code>if (m.containsKey(k))</code>	<code>if (m.count(k))</code> <code>if (m.contains(k)) (*)</code>
Check if the map is empty	<code>if (m.isEmpty())</code>	<code>if (m.empty())</code>
Retrieve or overwrite value associated with key k (auto-insert default if doesn't exist)	<code>int i = m[k];</code> <code>m[k] = i;</code>	<code>int i = m[k];</code> <code>m[k] = i;</code>

`std::map<K, V>`

stores a collection of

`std::pair<const K, V>`

*(I encourage you to think about why K is const.
What would happen if we could modify a key?)*

map as a collection of pair

We can iterate through the key-value pairs using a range based for loop

```
std::map<std::string, int> map;



for (auto kv : map) {
    // kv is a std::pair<const std::string, int>
    std::string key = kv.first;
    int value = kv.second;
}
```

map as a collection of pair

Structured bindings come in handy when iterating a map

```
std::map<std::string, int> map;  
  
for (const auto& [key, value] : map) {  
    // key has type const std::string&  
    // value has type const int&  
}
```

How is map implemented?



Binary Search Tree
(technically a *red-black tree*)

std::map<K, V> requires **K** to have an **operator<**

```
// ✅ OKAY - int has operator<
std::map<int, int> map1;
```

```
// ❌ ERROR - std::ifstream has no operator<
std::map<std::ifstream, int> map2;
```

std::set
#include <set>

std::set stores a collection of unique items

```
std::set<std::string> set {  
    "CS106L!",  
    "Keith",  
    "Sean",  
    "Nick",  
    "Chris"  
};
```



Stanford vs. STL set



What you want to do?	Stanford Set<char>	std::set<char>
Create an empty set	Set<char> s;	std::set<char> s;
Add k to the set	s.add(k);	s.insert(k);
Remove k from the set	s.remove(k);	s.erase(k);
Check if k is in the set (* C++20)	if (s.contains(k))	if (s.count(k)) if (s.contains(k)) (*)
Check if the set is empty	if (s.isEmpty())	if (s.empty())

std::set is an amoral std::map

std::set is an std::map without values

How is **set** implemented?

“Chris”
“CS106L”
“Keith”
“Nick”
“Sean”



Binary Search Tree
(technically a *red-black tree*)

std::unordered_map and std::unordered_set

```
#include <unordered_map>
#include <unordered_set>
```


`std::unordered_map`

- You can think of `unordered_map` as an optimized version of `map`
- It has the same interface as `map`

```
std::unordered_map<std::string, int> map {  
    { "Chris", 2 },  
    { "Nick", 51 },  
    { "Sean", 35 },  
};  
  
int sean = map["Sean"]; // 35  
map["Chris"] = 31;
```

How is `unordered_map` implemented?


- Remember, `map` is a collection of `std::pair`
- `unordered_map` stores a collection of n “buckets” of pairs



```
std::unordered_map  
<std::string, int> map {  
    { "Chris", 31 },  
    { "Nick", 51 },  
    { "Sean", 35 },  
};
```

How is `unordered_map` implemented?

- To add a key/value, we feed the key through a **hash function**
- The hash, modulo the bucket count, determines the pair's bucket no.





```
int x = map["CS106L"];
```

"CS106L" → $f(x)$ → 80489869

$$80489869 \bmod 5 = 4$$


What is a hash function?

- “Scrambles” a key into a `size_t` (64 bit)
- Small changes in the input should produce large changes in the output



How is `unordered_map` implemented?

- If two keys hash to the same bucket, we get a **hash collision**
- During lookup, we loop through bucket and check key equality
 - Two keys with the same hash are not necessarily equal!



`std::unordered_map<K, V>` requires `K` to have
a hash function (and equality)

Defined in header `<unordered_map>`

```
template<
    class Key,
    class T,
    class Hash = std::hash<Key>,
    class KeyEqual = std::equal_to<Key>,
    class Allocator = std::allocator<std::pair<const Key, T>>
> class unordered_map;
```

(We will learn more about this syntax later!)

`std::unordered_map<K, V>` requires K to be hashable

```
// ✓ OKAY - int is hashable
```

```
std::unordered_map<int, int> map1;
```


```
// ✗ ERROR - std::ifstream is not hashable
```

```
std::unordered_map<std::ifstream, int> map2;
```

Most basic types (`int`, `double`, `string`) are hashable by default

Why use `std::unordered_map`?

- **Load factor:** average number items per bucket
- `unordered_map` allows super fast lookup by keeping load factor small
- If load factor gets too large (above 1.0 by default), we **rehash**



Load Factor: 1.66

Load Factor: 1.0

Fun C++ Trivia: max_load_factor

You can control the max load factor before rehashing

```
std::unordered_map<std::string, int> map;

double lf = map.load_factor(); // Get current load factor
map.max_load_factor(2.0); // Set the max load factor

// Now the map will not rehash until load factor exceeds 2.0
// You should almost never need to do this,
// but it's a fun fact (good for parties!)
```

What makes a good hash function?

A good hash function minimizes the chance of a hash collision


```
// ✗ The worst possible hash

template <>
struct std::hash<MyType>
{
    std::size_t operator()(const MyType& k) const
    {
        return 0;
    }
};
```

(Don't worry too much about this syntax. We'll learn more later)

unordered_set is an **unordered_map** without values

```
std::unordered_set  
<std::string> set {  
    "Chris",  
    "Nick",  
    "Sean",  
    "CS106L"  
};
```



When to use `unordered_map` vs. `map`?

- `unordered_map` is *usually* faster than `map`
- However, it uses more memory (organized vs. disorganized garage)
- If your key type has no total order (`operator<`), use `unordered_map`!
- If you must choose, `unordered_map` is a safe bet

Summary of Data Structures

Space per Element ↓	ith element	Search	Insertion	Erase
	std::vector	Very Fast	Slow	Slow
	std::deque	Fast	Slow	Fast (front/back) Slow (all others)
	std::set	Slow	Fast	Fast
	std::map	Slow	Fast	Fast
	std::unordered_set	N/A	Very Fast	Very Fast
	std::unordered_map	N/A	Very Fast	Very Fast

Some more containers if you're curious!

std::array

A fixed-size array of items

std::list

A doubly linked list

std::multiset (+unordered)

A set that can contain duplicates

std::multimap (+unordered)

Can contain multiple values for the same key

Recap

- What the heck is the STL? What are templates?
 - "The Standard Template Library"
- Sequence Containers
 - A linear sequence of elements
 - `std::vector`, `std::deque`
- Associative Containers
 - A set of elements organized by unique keys
 - `std::map`, `std::set`, `std::unordered_map`, `std::unordered_set`