

Copying data is expensive

Quantifying the Energy Cost of Data Movement for Emerging Smart Phone Workloads on Mobile Platforms

Dhinakaran Pandiyan and Carole-Jean Wu

School of Computing, Informatics, and Decision Systems Engineering

Arizona State University

Tempe, Arizona 85281

Email: {dpandiya,carole-jean.wu}@asu.edu

moving data for a wide range of popular smart phone workloads. We find that a considerable amount of total device energy is spent in data movement (an average of 35% of the total device energy). Our results also indicate a relatively high stalled cycle

[\[source\]](#)

Copying data is expensive

Quantifying the Energy Cost of Data Movement in Scientific Applications

Gokcen Kestor*, Roberto Gioiosa*, Darren J. Kerbyson*, Adolfy Hoisie*

* Pacific Northwest National Laboratory

{gokcen.kestor, roberto.gioiosa, darren.kerbyson, adolfy.hoisie}@pnnl.gov

exascale systems. Projections show that the cost of moving data from memory is two orders of magnitudes higher than the cost of computing a double-precision register-to-register floating point operation. These

[\[source\]](#)

Last Time

- Special member functions handle the class lifecycle
 - Copy constructor `Type::Type(const Type& other);`
 - Copy assignment operator `Type& Type::operator=(const Type& other);`
 - Destructor `Type::~Type();`
- Compiler creates these for us
 - But... if we're managing memory, we need to override

Introducing... the Photo class

```
class Photo {  
public:  
    Photo(int width, int height);  
    Photo(const Photo& other);  
    Photo& operator=(const Photo& other);  
    ~Photo();  
private:  
    int width;  
    int height;  
    int* data;  
};
```

Photo Constructor

```
Photo::Photo(int width, int height)
: width(width)
, height(height)
, data(new int[width * height])
{}
```

Photo SMF: Copy Constructor

```
Photo::Photo(const Photo& other)
    : width(other.width)
    , height(other.height)
    , data(new int[width * height])
{
    std::copy(other.data, other.data + width * height, data);
}
```

Photo SMF: Copy Assignment

```
Photo &Photo::operator=(const Photo& other) {
    // Check for self assignment
    if (this == &other) return *this;

    delete[] data; // Clean up old pixels!

    // Copy over new pixels!
    width = other.width;
    height = other.height;
    data = new int[width * height];
    std::copy(other.data, other.data + width * height, data);
    return *this;
}
```

Photo SMF: Destructor

```
Photo::~Photo()
{
    delete[] data;
}
```

Your Turn

What special member functions get called at **(A)** and **(B)** below?

```
Photo takePhoto();  
  
int main() {  
    Photo coolestSelfie = takePhoto(); // (A) Copy Destruct RVO  
    Photo retake(0, 0);  
    retake = takePhoto();           // (B) Assign Destruct  
}
```

RVO = Return Value Optimization (compiler optimization)

Your Turn

What happened?!

```
Photo takePhoto();  
  
int main() {  
    Photo coolestSelfie = takePhoto(); // (A) Copy Destruct RVO  
    Photo retake(0, 0);  
    retake = takePhoto();           // (B) Assign Destruct  
}
```

The Problem



Photo

- width = 3840
- height = 2160
- data = 0x1024c3bd



retake = takePhoto();

The Problem



Photo

- width = 3840
- height = 2160
- data = 0x1024c3bd



Photo (retake)

- width = 0
- height = 0
- data = 0x153713f1

```
retake = takePhoto(); // Copy assignment
```

The Problem



Photo

- width = 3840
- height = 2160
- data = 0x1024c3bd



Photo (retake)

- width = 3840
- height = 2160
- data = 0x133210f1



```
retake = takePhoto(); // Copy assignment
```

The Problem



☠ Photo

- width = 3840
- height = 2160
- data = 0x1024c3bd



Photo (retake)

- width = 3840
- height = 2160
- data = 0x133210f1



```
retake = takePhoto(); // Destructor
```

What if we could reuse the memory instead?

The Solution: Move Semantics



Photo

- width = 3840
- height = 2160
- data = 0x1024c3bd



The Problem



Photo

- width = 3840
- height = 2160
- data = 0x1024c3bd



Photo (retake)

- width = 0
- height = 0
- data = 0x133210f1

```
retake = takePhoto(); // Move assignment
```

The Problem



retake = takePhoto(); // ~~Copy~~ Move assignment

The Problem



Photo

- width = 3840
- height = 2160
- data = 0x0

nullptr

Photo (retake)

- width = 3840
- height = 2160
- data = 0x1024c3bd



retake = takePhoto(); // ~~Copy~~ Move assignment

The Problem



✿ Photo

- width = 3840
- height = 2160
- data = 0x0

nullptr

Photo (retake)

- width = 3840
- height = 2160
- data = 0x1024c3bd



```
retake = takePhoto(); // Destructor
```

We **move** `takePhoto()` because it's temporary

Move vs. Copy Semantics

`takePhoto()` is temporary, so we can steal its resources!

Move vs. Copy Semantics

Is it always safe to move objects? Assume `get_pixel` accesses data

```
Photo takePhoto();  
  
void foo(Photo whoAmI) {  
    Photo selfie = whoAmI;          // What if we move here?  
    whoAmI.get_pixel(21, 24);       // ???  
}
```

What will happen
if we try to run
this code?

Move vs. Copy Semantics

- ✖ Since `selfie` stole `whoAmI`'s data, we end up dereferencing `nullptr`

```
Photo takePhoto();  
  
void foo(Photo whoAmI) {  
    Photo selfie = whoAmI;           // What if we move here?  
    whoAmI.get_pixel(21, 24);        // ✖ use-after-move  
}
```

Move vs. Copy Semantics

```
selfie1 = selfie2;  
// copy persistent objects (e.g. variables)
```

```
retake = takePhoto();  
// move temporary objects (e.g return values)
```

How does the compiler know whether to move or copy?

lvalues & rvalues

lvalues and rvalues generalize the idea of “temporariness” in C++

```
void foo(Photo pic) {  
    Photo beReal = pic;  
    Photo insta = takePhoto();  
}
```

pic is an lvalue!

takePhoto() is an
rvalue!

lvalues & rvalues

lvalues have a definite address, **rvalues** do not!

```
void foo(Photo pic) {  
    Photo* p1 = &pic; //   
    Photo* p2 = &takePhoto(); //  Doesn't work!  
}
```

pic is an lvalue!
We can take its
address!

takePhoto() is an
rvalue!
We **cannot** take its
address!

An **lvalue** can appear on either side of an **=**

```
x = y;
```

```
y = 5;
```

An **rvalue** can appear only right of an **=**

```
x = 5;
```

```
5 = y;
```

Your Turn

Which of the following right hand assignments are rvalues?

- Hint: which ones have a definite address?

int	a = 4;	rvalue
int&	b = a;	lvalue
vector<int>	c = {1, 2, 3};	rvalue
int	d = c[1];	lvalue
int*	e = &c[2];	rvalue
size_t	f = c.size();	rvalue
int	g = static_cast<int>(f);	lvalue

An **lvalue**'s lifetime is until the end of scope

An **rvalue**'s lifetime is until the end of line

Working towards move semantics

If we have an **lvalue**, how can we avoid copying its memory?

```
void uploadToInsta(Photo pic);

int main() {
    Photo selfie = takePhoto(); // selfie is lvalue
    uploadToInsta(selfie); // 🤔 Unnecessary copy is made here
}
```

Working towards move semantics

We can pass by reference! 

```
void uploadToInsta(Photo& pic);

int main() {
    Photo selfie = takePhoto(); // selfie is lvalue
    uploadToInsta(selfie); // ✅ No copy is made here
}
```

Working towards move semantics

- How can we avoid copying **rvalues**?
- What happens if we try to pass by reference?

```
void uploadToInsta(Photo& pic);  
  
int main() {  
    uploadToInsta(takePhoto()); // Does this work?  
}
```

✖ candidate function not viable: expects lvalue as 1st argument

lvalue reference

```
void upload(Photo& pic);  
  
int main() {  
    Photo selfie = takePhoto();  
    upload(selfie);  
}
```

rvalue reference

```
void upload(Photo&& pic);  
  
int main() {  
    upload(takePhoto());  
}
```

We can do whatever we want with **Photo&& pic**, it's temporary!

A few important points

- **Ivalue** references
 - Syntax: `Type&`
 - Persistent, must keep object in valid state after function terminates
- **rvalue** references
 - Syntax: `Type&&`
 - Temporary, we can steal (move) its resources
 - Object might end up in an invalid state, but that's okay! It's temporary!

```
//Hello. I want to take your Widget and play with it. It may be in a  
//different state than when you gave it to me, but it'll still be yours  
//when I'm finished. Trust me!  
void foo(Widget& w);
```

```
//Hello. Ooh, I like that Widget you have. You're not going to use it  
//anymore, are you? Please just give it to me. Thank you! It's my  
//responsibility now, so don't worry about it anymore, m'kay?  
void foo(Widget&& w);
```

[\[source\]](#)

**Key Idea: Overloading `&` and `&&` parameters
distinguish `lvalue` and `rvalue` references**

lvalue/rvalue overloading

```
void upload(Photo& pic);  
  
int main() {  
    Photo selfie = takePhoto();  
    upload(selfie);  
}
```

```
void upload(Photo&& pic);  
  
int main() {  
    upload(takePhoto());  
}
```

Compiler decides which version of **upload** to call depending on whether argument is **lvalue** or **rvalue**!

What we want!

```
Photo selfie1 = selfie2;  
// copy persistent objects (e.g. variables)  
  
retake = takePhoto();  
// move temporary objects (e.g return values)
```

Photo&

Photo&&

Let's overload the special member functions!

Copy constructor

```
Photo::Photo(const Photo& other)
: width(other.width)
, height(other.height)
, data(new int[width * height])
{
    std::copy(
        other.data,
        other.data + width * height,
        data
    );
}
```

Move constructor

```
Photo::Photo(Photo&& other)
: width(other.width)
, height(other.height)
{
    // other is temporary
    // Let's steal its
    // resources since we know
    // it's about to be gone!
}
```

```
Photo::Photo(Photo&& other)
: width(other.width)
, height(other.height)
, data(other.data)
{
    log() << "Photo(Photo&&)" << std::endl;
    other.data = nullptr;
}
```

Let's overload the special member functions!

Copy assignment operator

```
Photo& Photo::operator=(const Photo& other) {  
    if (this == &other) return *this;  
    delete[] data;  
    width = other.width;  
    height = other.height;  
    data = new int[width * height];  
    std::copy(other.data, other.data +  
              width * height, data);  
    return *this;  
}
```

Move assignment operator

```
Photo&  
Photo::operator=(Photo&& other)  
{  
    // other is temporary  
    // Let's steal its  
    // resources since we know  
    // it's about to be gone!  
}
```

```
Photo& Photo::operator=(Photo&& other)
{
    log() << "Photo::operator=(Photo&&)" << std::endl;

    // Clean up our data before assigning into this
    delete[] data;

    width = other.width;
    height = other.height;
    data = other.data;
    other.data = nullptr;

    return *this;
}
```

Forcing Move Semantics

- Usually, we let the compiler decide between `&` and `&&`
- Is that always the most efficient choice?
 - E.g. what if we know that an lvalue will never be used again?

Forcing Move Semantics

Line 3 *copies* each element into its new spot, even though the original value is never used again

```
1 void PhotoCollection::insert(const Photo& pic, int pos) {  
2     for (int i = size(); i > pos; i--)  
3         elems[i] = elems[i - 1]; // Shuffle elements down  
4     elems[i] = pic;  
5 }
```

Forcing Move Semantics

Solution: use move semantics

```
1 void PhotoCollection::insert(const Photo& pic, int pos) {  
2     for (int i = size(); i > pos; i--)  
3         elems[i] = std::move(elems[i - 1]);  
4     elems[i] = pic;  
5 }
```

`std::move` casts an `lvalue` to an `rvalue`, allowing compiler to select correct overload

Be wary of `std::move`

If we move an lvalue, what happens to it afterwards?

```
Photo takePhoto();  
  
void foo(Photo whoAmI)  
{  
    Photo selfie = std::move(whoAmI);  
    whoAmI.get_pixel(21, 24); // ???  
}
```

✖ If we move, `whoAmI` ends up in an unknown state!

Use `std::move` to implement move operations!

```
class Photo {  
public:  
    Photo::Photo(Photo&& other) {  
        keywords = other.keywords;  
    }  
  
private:  
    std::vector<string> keywords;  
};
```

We know that `other` is temporary! So do we *really* need to make a copy of `other.keywords`?

Use `std::move` to implement move operations!

```
class Photo {  
public:  
    Photo::Photo(Photo&& other) {  
        keywords = std::move(other.keywords);  
    }  
private:  
    std::vector<string> keywords;  
};
```

Solution: force move semantics by using `std::move`

std::move doesn't do anything special!

- std::move just type casts an lvalue to an rvalue

Return value

```
static_cast<typename std::remove_reference<T>::type&&>(t)
```

- Like const_cast, we "opt in" to potentially error-prone behaviour
 - What if we try to use an object after it's been moved! 🚨 SOS 🚨
 - Try to avoid explicitly using std::move unless you have good reason!
 - E.g. performance really matters, you know for sure the object won't be used!

So many SMFs... 🥺
Do I need to define them all!?



No! Rule of
zero, three,
and five!

Rule of Zero

- If a class doesn't manage memory (or another external resource), the compiler generated versions of the SMFs are sufficient!
- **Example:** Compiler generated SMFs of `Post` will call SMFs of `Photo` and `std::string`

```
struct Post {  
    Photo photo;  
    std::string user;  
};
```

Rule of Three

- If a class manages external resources, we must define **copy assignment/constructor**
- If we don't, compiler-generated SMF won't copy underlying resource
 - This will lead to bugs, e.g. two **Photo**'s referring to the same underlying data

Rule of Three: If you need any one of these, you need them all:

- Destructor
- Copy Assignment
- Copy Constructor

Rule of Five

- If we defined **copy constructor/assignment** and **destructor**, we should also define **move constructor/assignment**
- This is not required, but our code will be slower as it involves unnecessary copying

Rule of Five: If you need any of these, you probably want them all:

- Destructor
- Copy Assignment
- Copy Constructor
- Move Assignment (Optional)
- Move Constructor (Optional)

What we covered

- SMFs Recap
 - SMFs manage the lifetime of a class instance.
- The Problem
 - The copy SMFs make a copy of every object, including temporary ones!
- lvalues and rvalues
 - An lvalue is a persistent object. An rvalue is a temporary one.
- Move Semantics
 - We can define “move” overloads of the copy constructor/assignment operator
- **std::move** and SMFs
 - std::move opts into move semantics, but be careful! Rule of zero/three/five